

Chapter 11 Notes
Threads
Advanced Programming in the Unix Environment 3e
MCIT 5950

Tyler Ryan

February 15, 2023

Contents

1	Before Reading	2
2	High Level Overview	2
3	Thread Concepts	2
4	Basic Function Prototypes	3
4.1	Comparing Process Functions and Thread Functions	4
5	Thread Identification	4
6	Thread Creation	4
6.1	Code: Printing Thread IDs	4
7	Thread Termination	6
7.1	Code: Getting Exit Codes From Threads	6
7.2	Code: Incorrectly Passing Values Between Threads	8
7.2.1	Code Challenge 1: Returning Multiple Values From Thread	10
8	Thread Synchronization	11
8.1	Function Prototypes	11
8.2	Mutexes	11
8.2.1	Code: Basic Mutex Implementation	11
8.2.2	Code Challenge 2: Dual Threaded Mutex	14
8.3	Thread Deadlocking	14
8.3.1	Code: Locking Multiple Mutexes	14
8.4	Reader-Writer Locks	20
8.4.1	Function Prototypes	20
8.5	Condition Variables	21
8.5.1	Function Prototypes	21
9	Code Challenge Solutions	22
9.1	Challenge 1	22
9.2	Challenge 2	24

1 Before Reading

Prerequisite Knowledge

You should know a little about the following subjects:

- Generic Types using Void Pointers
- Single Linked List
- Double Linked List
- Queues
- Hash Tables (Using Linked Lists)

Installing Missing Man Pages

Run *dnf install man-pages-posix* to install all **pthread.h** manual pages.

2 High Level Overview

- See how threads of control (i.e. Threads) can be used to perform multiple tasks within the environment of a single process.
- All threads within a single process have access to (share) the **same** process components, such as file descriptors and memory.
 - Compare this to this to multiples processes.
- Threads and POSIX.1 primitives available to create and destroy them.
- Thread synchronization problem.
- 3 fundamental synchronization mechanisms:
 - Mutexes
 - Reader-Writer Locks
 - Conditional Variables

3 Thread Concepts

- With multiple threads (of control), a program can be designed to do more than on thing at a time **within a single process**, with each thread handling separate tasks.
- Benefits of Threads:
 - Code that deals with asynchronous events can be simplified by assigning a separate thread to handle each event type.
 - * Each thread can then handle its event using a synchronous programming model.
 - * A synchronous programming model is much simpler than an asynchronous one.
 - Unlike multiple processes which have to use complex mechanisms provided by the operating system to share memory and file descriptors, **Threads automatically have access to the same memory address space and file descriptors.**
 - Unlike processes which implicitly serializes multiple tasks, threads can interleave indepentent tasks by assigning a separate thread per task.
 - * Two tasks can be interleaved only if they don't depend on the processing performed by each other.
 - Interactive programs can realize **improved response time** by using multiple threads to separate the porions of the program that deal with user input and output from the other parts of the program.

- The benefits of a multithreaded programming model can be realized even if your program is running on a uniprocessor.
- A thread consists of the information necessary to represent an **execution context** within a process.
 - Thread ID
 - Set of Registers
 - A Stack
 - A Scheduling priority and policy
 - A Signal Mask
 - An Errno Variable
 - Thread-Specific Data
- Everything within a process is sharable among the threads in that process including:
 - The text of the executable program
 - The program's global and heap memory
 - The Stacks
 - The File Descriptors

4 Basic Function Prototypes

```
#include <pthread.h>
// Data Type
pthread_t

// Compare the value of 2 threads
int pthread_equal(pthread_t tid1, pthread_t tid2);

// Get current thread's ID
pthread_t pthread_self(void);

// Create a thread
int pthread_create(pthread_t * restrict tidp,
                  const pthread_attr_t * restrict attr,
                  void * (*start_rtn)(void *),
                  void * restrict arg);

// Returns from a thread
void pthread_exit(void * rval_ptr);

// Like wait() for threads
int pthread_join(pthread_t thread, void **rval_ptr);

// Like sending pthread_exit() to another thread
int pthread_cancel(pthread_t tid);

// Cleanup functions just before threads exit
void pthread_cleanup_push(void (*rtn)(void *), void * arg);
// "int execute" is more like "bool execute"
void pthread_cleanup_pop(int execute);
```

4.1 Comparing Process Functions and Thread Functions

Table 1: Process Functions vs Thread Functions

Process Primitive	Thread Primitive	Description
fork	pthread_create	Create a new flow of control
exit	pthread_exit	Exit from an existing flow control
waitpid	pthread_join	Get exit status from flow control
atexit	pthread_cleanup_push	Register function to be called at exit from flow of control
getpid	pthread_self	Get ID for flow of control
abort	pthread_cancel	Request abnormal termination of flow of control
kill	pthread_kill	Send signals to other flows of control

5 Thread Identification

- Every thread has a Thread ID
- These Thread IDs are only unique inside each process.
- Each thread shares the same Process ID (PID)

6 Thread Creation

- All programs start as single threaded processes.
- When a thread is created, just like with processes, there is no guarantee which runs first: The newly created thread or the calling thread (program).
 - Has access to the process address space.
 - Inherits the calling thread's floating-point environment and signal mask.
 - However, the set of *pending signals* for the thread is cleared.
- If you need to pass more than one argument to the **start_rtn** function, then you need to store them in a struct and pass the *address* of that struct to the **arg** parameter.
- pthread functions usually return an error code when they fail. they don't set **errno** like the other POSIX functions.

6.1 Code: Printing Thread IDs

```
#include <stdio.h>
#include <stdlib.h>
#include <unistd.h>
#include <pthread.h>
#include <sys/types.h>
#include <string.h>

/* Write a function that prints the current thread id
 * */

pthread_t ntid;

/* Notice how similar the Hex numbers (address ranges) are
 * This means that they are near each other in memory.
 *
 * Running on Linux (Fedora 37), these are the results.
 *
 * PIDs are the same.
 *
```

```

    * TID addresses are different ,
    * First 4 positions are the same.
    * Last 5 positions are the same.
    */
void
print_tid(const char * s)
{
    pthread_t tid = pthread_self();
    printf("%s (TID): %u, (HEX) %08x%x, (PID) %u\n",
           s,
           (unsigned int) tid,
           (unsigned int) tid,
           (unsigned int) getpid());
}

/* When you create a new thread, you must pass in a function.
   * In our case, our function just calls our print_tid() function.
   */
void *
thread_fn(void * arg)
{
    print_tid("new thread: ");

    // Note: In Linux, returning just 0 is allowable.
    return (void *) 0;
}

int
main(void)
{
    int      err;

    // This will call a thread that will then print it's TID.
    err = pthread_create(&tid, NULL, thread_fn, NULL);

    if (err != 0)
    {
        printf("error: can't create thread: %d\n", strerror(err));
        exit(EXIT_FAILURE);
    }

    // Print the main thread's TID
    print_tid("main thread: ");
    sleep(1);
    exit(EXIT_SUCCESS);
}

```

Output

```

main thread: (TID): 1975252800, (HEX) 75bbf7401fef9, (PID) 206267015
new thread: (TID): 1975248576, (HEX) 75bbe6c01fef9, (PID) 206267015

```

7 Thread Termination

- A single thread can exit in 3 ways, stopping its flow of control without terminating the entire process.
 1. The thread can return from the start routine.
 - The return value is the thread's exit code
 2. The thread can be canceled by another thread in the **same process**.
 3. The thread can call **pthread_exit()**
- The thread that calls **pthread_join()** will block until the specified thread calls **pthread_exit()** returns from its start routine or is canceled.
 - If the thread returns from its routine, **rval_ptr** will contain the return code.
 - If the thread was canceled, the memory location specified by **rval_ptr** is set to **PTHREAD_CANCELED**.
- By calling **pthread_join()**, the joining thread is automatically placed in a detached state so that its resources can be recovered.

7.1 Code: Getting Exit Codes From Threads

```
#include <stdio.h>
#include <stdlib.h>
#include <unistd.h>
#include <pthread.h>
#include <string.h>

/* A simple program that shows how to
 * obtain the exit status and the
 * return value of a thread using either
 * "return" or "pthread_exit()"
 */

void *
thread_func1(void * arg)
{
    printf("thread 1 returning\n");
    return ((void*)1);
}

// See if you can return a char *
// I'd prefer this one
void *
thread_func2(void * arg)
{
    printf("thread 2 exiting\n");
    // To return a char *
    // pthread_exit("HELLER\n");
    pthread_exit((void*)2);
}
```

```
int
main(void)
{
    pthread_t   thread1, thread2;
    int         err;
    void        * thread_return;

    // Create thread 1, which returns 1
    err = pthread_create(&thread1, NULL, thread_func1, NULL);
    if (err !=0)
    {
        printf("error: can't create thread 1: %s\n", strerror(err));
        exit(EXIT_FAILURE);
    }

    // Create thread 2, which returns 2 using pthread_exit()
    err = pthread_create(&thread2, NULL, thread_func2, NULL);
    if (err != 0)
    {
        printf("error: can't create thread 2: %s\n", strerror(err));
        exit(EXIT_FAILURE);
    }

    /* Waits for thread 1 and returns the exit status to err,
     * and your whatever thread_func1() returned to thread_return.
     */
    err = pthread_join(thread1, &thread_return);
    // err == 0
    // thread_return == 1

    if (err !=0)
    {
        printf("error: can't join with thread 1: %s\n",
               strerror(err));
        exit(EXIT_FAILURE);
    }

    printf("thread 1 exit code: %d\n", (int*) thread_return);

    /* Waits for thread 2 and returns the exit status to err,
     * and your whatever thread_func2() returned to thread_return.
     */
    err = pthread_join(thread2, &thread_return);
    // err == 0
    // thread_return == 2

    if (err !=0)
    {
        printf("error: can't join with thread 2: %s\n",
               strerror(err));
    }
}
```

```

        exit(EXIT_FAILURE);
    }

    printf("thread 2 exit code: %d\n", (int*) thread_return);
    // To return a char *
    //printf("thread 2 exit code: %s\n", (char*) thread_return);

    exit(EXIT_SUCCESS);
}

```

Output

```

thread 2 exiting
thread 1 returning
thread 1 exit code: 1
thread 2 exit code: 2

```

- Be careful when passing values from thread to thread.
- For example, if you create a struct on one thread's stack and try to pass it to another, it will return random values as that struct was destroyed when that thread was destroyed.
- In essence, when values create within threads die with their threads unless made static or put on the heap.

7.2 Code: Incorrectly Passing Values Between Threads

```

#include <stdio.h>
#include <stdlib.h>
#include <unistd.h>
#include <pthread.h>

struct foo { int a, b, c, d; } ;

void
print_foo(const char *s, const struct foo *fp)
{
    printf(s);
    printf(" structure at 0x%x\n", *(unsigned int*)fp);
    printf(" foo.a = %d\n", fp->a);
    printf(" foo.b = %d\n", fp->b);
    printf(" foo.c = %d\n", fp->c);
    printf(" foo.d = %d\n", fp->d);
}

void *
thread_func1(void * arg)
{
    // Creating a foo struct on this thread's stack
    struct foo foo = {1,2,3,4};
    print_foo("thread 1:\n", &foo);
}

```



```
    /* Because foo was not made static or put on the
     * heap, it will die when we exit this function.
     */
    pthread_exit((void*)&foo);
}

void *
thread_func2(void * arg)
{
    printf("thread 2: ID is %ld\n", pthread_self());
    pthread_exit((void*)0);
}

int
main(void)
{
    /* For the sake of simplicity, I'm not handling
     * any errors
     */
    pthread_t tid1, tid2;
    struct foo *fp;
    int err;

    // CREATES a foo struct WITHIN this thread
    err = pthread_create(&tid1, NULL, thread_func1, NULL);

    /* Returns a pointer to the foo struct here.
     * However, that thread has already been terminated.
     * If you tried to access the value of fp which points
     * to memory that has been terminated, you will get
     * junk.
     */
    err = pthread_join(tid1, (void*)&fp);

    sleep(1);

    printf("parent starting second thread\n");

    err = pthread_create(&tid2, NULL, thread_func2, NULL);

    sleep(1);
    print_foo("parent:\n", fp);

    exit(EXIT_SUCCESS);
}
```

Output

```
thread 1:
  structure at 0x1
  foo.a = 1
  foo.b = 2
  foo.c = 3
  foo.d = 4
parent starting second thread
thread 2: ID is 140514740258496
parent:
  structure at 0x232dc6c0
  foo.a = 590202560
  foo.b = 32716
  foo.c = 0
  foo.d = 0
```

The parent's output *should* be the same as the first thread, however it isn't because its data was destroyed when the first thread exited.

7.2.1 Code Challenge 1: Returning Multiple Values From Thread

Try to get this to work. Hint: Use global/heap space.

Expected Output

```
thread 1:
  structure at 0x1
  foo.a = 1
  foo.b = 2
  foo.c = 3
  foo.d = 4
parent starting second thread
thread 2: ID is 140436664235712
parent:
  structure at 0x1
  foo.a = 1
  foo.b = 2
  foo.c = 3
  foo.d = 4
```

See Listing [2](#) for a **solution**.

8 Thread Synchronization

8.1 Function Prototypes

```
#include <pthread.h>

// ----- MUTEX Functions -----
// Mutex Data Type
pthread_mutex_t

// Use for statically-allocated Mutexes
PTHREAD_MUTEX_INITIALIZER

int pthread_mutex_init(pthread_mutex_t * restrict mutex
                      const pthread_mutexattr_t * restrict attr);
int pthread_mutex_destroy(pthread_mutex_t * mutex);

int pthread_mutex_lock(pthread_mutex_t * mutex);
int pthread_mutex_trylock(pthread_mutex_t * mutex);
int pthread_mutex_unlock(pthread_mutex_t * mutex);

                                Return: 0 if OK, error number on failure
// -----
```

8.2 Mutexes

A Mutex variable is represented by **pthread_mutex_t**

Use **PTHREAD_MUTEX_INITIALIZER** for statically-allocated mutexes.

8.2.1 Code: Basic Mutex Implementation

```
#include <stdio.h>
#include <stdlib.h>
#include <unistd.h>
#include <pthread.h>

/* Using a single thread, get this
 * code to work.
 */

struct foo
{
    int          f_count;
    pthread_mutex_t f_lock;
};

struct foo *
foo_alloc(void)
{
```

```

    struct foo * fp;
    fp = malloc(sizeof(struct foo));
    // Keep track of f_lock
    int res = pthread_mutex_init(&fp->f_lock , NULL);

    if (res != 0)
    {
        free(fp);
        return (NULL);
    }

    return fp;
};

void
foo_hold(struct foo * fp)
{
    // Lock f_lock
    pthread_mutex_lock(&fp->f_lock);
    // Do something
    fp->f_count++;
    // Unlock f_lock
    pthread_mutex_unlock(&fp->f_lock);
}

void
foo_release(struct foo * fp)
{
    /* "—fp->f_count" is interesting code.
     *
     * What it does is pre-decrement the actual
     * f_count amount by 1.
     *
     * if fp->f_count was 1, then —fp->f_count
     * would == 0.
     *
     * If that is the case, then we can go ahead and
     * do the following:
     * 1. Unlock it
     * 2. Destroy it (make it uninitialized)
     * 3. Free it's memory
     */

    if (--fp->f_count == 0)
    {
        pthread_mutex_unlock(&fp->f_lock);
        pthread_mutex_destroy(&fp->f_lock);
        free(fp);
        fp = NULL;
    }
}

```

```

    /* However, if there are multiple foo objects,
     * then just unlock this particular one.
     */
    else
    {
        pthread_mutex_unlock(&fp->f_lock);
    }
}

int
main(void)
{
    /* Note: For the sake of simplicity, this example
     * is only showing how to use these functions with
     * 1 thread.
     */
    struct foo * fpointer = foo_alloc();

    // Get 200
    for (int i = 0; i < 200; i++)
    {
        foo_hold(fpointer);
    }

    printf("200: %d\n", fpointner->f_count); // 200

    // Release all but 1
    for (int i = 0; i < 199; i++)
    {
        foo_release(fpointer);
    }

    printf("1: %d\n", fpointer->f_count); // 1

    // Release the last one
    foo_release(fpointer);
    // Invalid memory produces junk information
    printf("Destroyed: %d\n", fpointer->f_count);

    exit(EXIT_SUCCESS);
}

```

Output

```

200: 200
1: 1
Destroyed: 6846

```

This example shows how to use a mutex with only a single thread.

8.2.2 Code Challenge 2: Dual Threaded Mutex

Try running 2 threads. Each thread should increment `fpinter` by 1 up to 100. So the total should still end up being 200 (i.e. the same output as above).

See Listing 3 for a **solution**.

8.3 Thread Deadlocking

- A thread will deadlock itself if it tries to lock the same mutex twice.
- If more than one mutex is used in a program a deadlock can occur if the thread is allowed to hold a mutex and block while trying to lock a second mutex at the same time the other thread holding the second mutex tries to lock the first mutex.
- In this case, neither thread can proceed, because each needs a resource that is held by the other, thus the deadlock.
- To avoid deadlocks, ensure that when there is a need to acquire two mutexes at the same time, to always lock them in the same order.

8.3.1 Code: Locking Multiple Mutexes

Note:

- At this point the textbook gets pretty wild with its examples.
- The code in the book wasn't a working example. I've added a few things and made it work using a single thread for simplicity.

In this example, deadlocks are avoided by ensuring that when two mutexes need to be acquired at the same time, they are a locked in the same order.

- The second mutex protects a hash list that is used to keep track of the **foo** data structures.
- The **hashlock** mutex protects both the **fh** hash table and the **f_next** hash link field in the **foo** structure.
- The **f_lock** mutex in the **foo** structure protects the access to the remainder of the **foo** structure's fields.

Listing 1: Multiple Mutex Locking with Hash Table

```
#include <stdio.h>
#include <stdlib.h>
#include <pthread.h>
/* Please note, the book did not give a
 * working example of this program.
 */

// Number of slots
#define NHASH 12
/* HASH(fp) basically reduces our hash number to fit
 * within our NHASH slots
 *
 * Ex.
 * int hash = 5000;
```

```

    * int smaller_hash = HASH(hash); // smaller_hash = 8;
    */
#define HASH(fp) (((unsigned long)fp) % NHASH)

// Struct we are trying to protect
struct foo
{
    int          f_count;
    pthread_mutex_t f_lock;
    struct foo    * f_next;
    int          f_id;
    /* ... more stuff here ... */
};

// Functions
struct foo * foo_alloc(int);
void    foo_hold(struct foo * fp);
struct foo * foo_find(int id);
struct foo * my_find(struct foo * entry);
void    foo_release(struct foo * fp);
void    print_ht_values(void);

// Global variable
// An array of foo
struct foo * fh[NHASH];
// Our global lock/unlock variable
pthread_mutex_t hashlock = PTHREAD_MUTEX_INITIALIZER;

int
main(void)
{
    /* TODO: Mess around with some stuff you
     * are unsure about here.
     */
    struct foo * a = foo_alloc(150);
    struct foo * b = foo_alloc(200);
    struct foo * c = foo_alloc(300);

    // Note that these both will stay the same throughout the
    // same execution, but can/will change when reran.
    printf("Actual Value of fp: %ld\n", a);
    printf("Hash value: %ld\n", HASH(a));

    print_ht_values();

    int a_val = foo_find(150)->f_id;
    printf("Just found %d by searching\n", a_val);

    // Uses the hash feature of the hash table to find values

```

```

    int my_val = my_find(a)->f_id;
    printf("I just found %d\n", my_val);

    // This doesn't seem to work
    foo_release(a);
    printf("Removed 150\n");
    printf("New Table\n");
    print_ht_values();
}

void print_ht_values(void)
{
    // Loop through the array
    for (int i = 0; i < NHASH; i++)
    {
        long int fp_value = (long int) fh[i];
        int hash_value = HASH(fp_value);
        // Just print the ones that were allocated space
        if (fp_value != 0)
        {
            printf("index %d, fp: 0x%x, val: %d\n",
                    hash_value,
                    (unsigned int) fp_value,
                    fh[i]->f_id
            );
        }
    }
}

struct foo *
foo_alloc(int id)
{
    struct foo * fp;
    int ht_index;

    fp = malloc(sizeof(struct foo));
    // If malloc worked
    if (fp != NULL)
    {
        // Set our count to 1
        fp->f_count = 1;

        // If this fails, free fp
        if (pthread_mutex_init(&fp->f_lock, NULL) != 0)
        {
            free(fp);
            return NULL;
        }
        /* fp is just a number we're trying to squeeze down
         * below our NHASH number.

```



```

        * This will create the index where we will store
        * our values
        */
    ht_index = HASH(fp);
    // Lock down the struct so we can mess with it safely
    pthread_mutex_lock(&hashlock);
    // point next node to this index
    fp->f_next = fh[ht_index];
    // Set value of that index to this fp
    fh[ht_index] = fp;
    fp->f_id = id;

    // Never allowed to unlock this node?
    pthread_mutex_lock(&fp->f_lock);
    // Unlock struct
    pthread_mutex_unlock(&hashlock);
    /* ... continue initializing ...*/
}
return fp;
}

void
foo_hold(struct foo * fp)
{
    pthread_mutex_lock(&hashlock);
    fp->f_count++;
    pthread_mutex_unlock(&hashlock);
}

// this doesn't utilize the hash feature when finding
// values in the hash table lol
// It just loops through everything.
struct foo *
foo_find(int id)
{
    struct foo * fp;
    int          ht_index;

    pthread_mutex_lock(&hashlock);
    // For each index of the array
    for (ht_index = 0; ht_index < NHASH; ht_index++)
    {
        // for each node at each index
        for (fp = fh[ht_index]; fp != NULL; fp = fp->f_next)
        {
            // If the node's id matches our id
            if (fp->f_id == id)
            {
                // increase fp's memory location by 1?
                fp->f_count++;
            }
        }
    }
}

```

```

        pthread_mutex_unlock(&hashlock);
        return fp;
    }
}

// ID couldn't be found
pthread_mutex_unlock(&hashlock);
return NULL;
}

/* This way uses the hash feature.
 * Notice how I've reduced the nested for loop
 * to just a single for loop.
 *
 * If you were to assert that all keys were unique,
 * then you could create a hash function where you
 * could look up the actual value (instead of the object)
 * like you can with the book's function.
 */
struct foo *
my_find(struct foo * entry)
{
    struct foo * fp;
    int          ht_index;

    pthread_mutex_lock(&hashlock);

    unsigned int hash_code = HASH(entry);
    printf("Hash Code %d\n", hash_code);

    for (fp = fh[hash_code]; fp != NULL; fp = fp->f_next)
    {
        // If the node's id matches our id
        if (fp->f_id == entry->f_id)
        {
            // increase fp's memory location by 1?
            fp->f_count++;
            pthread_mutex_unlock(&hashlock);
            return fp;
        }
    }
    // ID couldn't be found
    pthread_mutex_unlock(&hashlock);
    return NULL;
}

// This function DOES utilize the hash feature
void
foo_release(struct foo * fp)

```

```

{
    struct foo    *temp_foo_ptr;
    int           ht_index;

    pthread_mutex_lock(&hashlock);

    // If this is the last "instance"
    if (--fp->f_count == 0)
    {
        ht_index      = HASH(fp);
        temp_foo_ptr = fh[ht_index];

        // If temp_fp happens to be the first node
        if (temp_foo_ptr == fp)
        {
            // Hashtable[index] = this node
            fh[ht_index] = fp->f_next;
        }
        // else loop through the nodes, until
        // the correct node is found.
        else
        {
            while (temp_foo_ptr->f_next != fp)
            {
                temp_foo_ptr = temp_foo_ptr->f_next;
            }
            // At this point, if it's there
            // this will set us there.
            temp_foo_ptr->f_next = fp->f_next;
        }

        // Remove that one from the table
        // 1. Unlock the table
        // 2. Destroy the node's lock
        // 3. Free from memory
        pthread_mutex_unlock(&hashlock);
        pthread_mutex_destroy(&fp->f_lock);
        free(fp);
    }
    // If it wasn't found, unlock table and do nothing.
    else
    {
        pthread_mutex_unlock(&hashlock);
    }
}

```

8.4 Reader-Writer Locks

8.4.1 Function Prototypes

```
#include <pthread.h>

// Initialize
int pthread_rwlock_init(pthread_rwlock_t * restrict rwlock,
                        const pthread_rwlockattr_t * restrict attr);

// Destroy/Clean Up
int pthread_rwlock_destroy(pthread_rwlock_t * rwlock);

// Read Lock
int pthread_rwlock_rdlock(pthread_rwlock_t * rwlock);
// Write Lock
int pthread_rwlock_wrlock(pthread_rwlock_t * rwlock);
// Unlock
int pthread_rwlock_unlock(pthread_rwlock_t * rwlock);

// Try Read
int pthread_rwlock_tryrdlock(pthread_rwlock_t * rwlock);
// Try Write
int pthread_rwlock_trywrlock(pthread_rwlock_t * rwlock);
```

- Also called **shared-exclusive locks**.
 - Read mode == shared mode.
 - Write mode == exclusive mode.
- Reader-writer locks are similar to mutexes, except that they allow for higher degrees of parallelism.
- With a mutex, the state is either locked or unlocked, and only one thread can lock it at a time.
- Only one thread at a time can hold a reader-writer lock in write mode, but multiple threads can hold a reader-writer lock in read mode at the same time.
- Reader-writer locks are well suited for situations in which data structures are read more often than they are modified.
- When a reader-writer lock is write-locked, all threads attempting to lock it block until it is unlocked.
- When a reader-writer lock is read-locked, all threads attempting to lock it in read mode are given access, but any threads attempting to lock it in write mode block until all the threads have relinquished their read locks.
 - Implementations vary, but reader-writer locks usually block additional readers if a lock is already held in read mode and a thread is blocked trying to acquire the lock in write mode.
- When a reader-writer lock is held in write mode, the data structure it protects can be modified safely, since only one thread at a time can hold the lock in write mode.
 - Prevents constant stream of readers from starving waiting writers.

8.5 Condition Variables

8.5.1 Function Prototypes

```
#include <pthread.h>

// Data type
pthread_cond_t

// For statically-allocated condition variables.
PTHREAD_COND_INITIALIZER

int pthread_cond_init(pthread_cond_t * restrict cond,
                     pthread_condattr_t * restrict attr);
int pthread_cond_destroy(pthread_cond_t * cond);

int pthread_cond_wait(pthread_cond_t * restrict cond,
                     pthread_mutex_t * restrict mutex);
int pthread_cond_timedwait(pthread_cond_t * restrict cond,
                           pthread_mutex_t * restrict mutex,
                           const struct timespec * restrict timeout);
// timespec structure used in ...timedwait()
struct timespec
{
    time_t  tv_sec;    // Seconds
    long    tv_nsec;   // Nanoseconds
}

// Wakes up 1 thread waiting on a condition
int pthread_cond_signal(pthread_cond_t * cond);
// Wakes up all threads waiting on a condition
int pthread_cond_broadcast(pthread_cond_t * cond):
```

- Condition variables are another synchronization mechanism available to threads.
- When used with mutexes, condition variables allow threads to wait in a race-free way for arbitrary conditions to occur.
- The condition itself is protected by a mutex.
- A thread must first lock the mutex to change the condition state.
 - Other threads will not notice the change until they acquire the mutex.
 - As the mutex must be locked to be able to evaluate the condition.

9 Code Challenge Solutions

These are just some problems I came up with.

They might or might not:

- Help
- Be correct
- Be worth your time

9.1 Challenge 1

Here I used **malloc** to solve this problem, but you could've also used **static**.

Listing 2: Solution: Returning Multiple Values From A Thread

```
#include <stdio.h>
#include <stdlib.h>
#include <unistd.h>
#include <pthread.h>
#include <string.h>

/* Write a program that returns a struct
 * (i.e. multiple values) from a thread
 * using malloc
 */

struct foo { int a, b, c, d; } ;

void
print_foo(const char *s, const struct foo *fp)
{
    printf(s);
    printf(" structure at 0x%x\n", *(unsigned int*)fp);
    printf(" foo.a = %d\n", fp->a);
    printf(" foo.b = %d\n", fp->b);
    printf(" foo.c = %d\n", fp->c);
    printf(" foo.d = %d\n", fp->d);
}

// Create a function struct within the thread and returns it
void *
thread_func1(void * arg)
{
    // If you don't put foo on the heap, it will die
    // when this thread dies.
    struct foo * foo = (struct foo*) malloc(sizeof(struct foo));
    foo->a = 1;
    foo->b = 2;
    foo->c = 3;
    foo->d = 4;
```

```
// Proof foo exists within the thread
print_foo("thread 1:\n", foo);
pthread_exit((void *)foo);
}

void *
thread_func2(void * arg)
{
    printf("thread 2: ID is %ld\n", pthread_self());
    pthread_exit((void *)1);
}

void print_error(char * msg, int err)
{
    printf("%s\n", strerror(err));
    exit(EXIT_FAILURE);
}

int
main(void)
{
    int      err;
    struct   foo *fp;
    pthread_t tid1, tid2;

    // Creates a foo struct within this thread
    err = pthread_create(&tid1, NULL, thread_func1, NULL);
    if (err != 0) print_error("", err);

    // Returns a pointer to the foo struct here
    err = pthread_join(tid1, (void*)&fp);
    if (err != 0) print_error("", err);

    sleep(1);

    printf("parent starting second thread\n");

    err = pthread_create(&tid2, NULL, thread_func2, NULL);
    if (err != 0) print_error("", err);

    sleep(1);
    print_foo("parent:\n", fp);

    // You must free when you use malloc
    free(fp);

    exit(EXIT_SUCCESS);
}
```

9.2 Challenge 2

Listing 3: Solution: Dual Threaded Mutex

```

#include <stdio.h>
#include <stdlib.h>
#include <unistd.h>
#include <pthread.h>

/* Using multiple threads ,
 * get fp->f_count to equal to 200.
 */

struct foo
{
    int          f_count;
    pthread_mutex_t  f_lock;
};

struct foo *
foo_alloc(void)
{
    struct foo * fp;
    fp = malloc(sizeof(struct foo));
    // Keep track of f_lock
    int res = pthread_mutex_init(&fp->f_lock , NULL);

    if (res != 0)
    {
        free(fp);
        return (NULL);
    }

    return fp;
};

void
foo_hold(struct foo * fp)
{
    // Lock f_lock
    pthread_mutex_lock(&fp->f_lock);
    // Do something
    fp->f_count++;
    // Unlock f_lock
    pthread_mutex_unlock(&fp->f_lock);
}

void
foo_release(struct foo * fp)
{
    /* "—fp->f_count" is interesting code.
     *

```



```

    * What it does is pre-decrement the actual
    * f_count amount by 1.
    *
    * if fp->f_count was 1, then --fp->f_count
    * would == 0
    *
    * If that is the case, then we can go ahead and
    * do the following:
    * 1. Unlock it
    * 2. Destroy it (make it uninitialized)
    * 3. Free it's memory
    */
    if (--fp->f_count == 0)
    {
        pthread_mutex_unlock(&fp->f_lock);
        pthread_mutex_destroy(&fp->f_lock);
        free(fp);
    }

    else
    {
        pthread_mutex_unlock(&fp->f_lock);
    }
}

void
empty_foo(struct foo * fp)
{
    int total = fp->f_count;

    for (int i = 0; i < total; i++)
    {
        foo_release(fp);
    }
}

void *
thread_func(void * arg)
{
    for (int i = 0; i < 100; i++)
    {
        foo_hold(arg);
    }
}

int
main(void)
{
    struct foo * fp = foo_alloc();
    int status;

```

```
pthread_t tid1, tid2;

// returns 0 upon success. I'm just not catching
// these errors for simplicity's sake.
status = pthread_create(&tid1, NULL, thread_func, fp);
status = pthread_create(&tid2, NULL, thread_func, fp);

if (pthread_join(tid1, NULL) != 0)
{
    printf("error: thread 1\n");
    exit(EXIT_FAILURE);
}

// Same thing as above with different syntax
if (pthread_join(tid2, NULL))
{
    printf("error: thread 2\n");
    exit(EXIT_FAILURE);
}

printf("200: %d\n", fp->f_count);

// Unlock, Destroy, and Free structure memory
empty_foo(fp);
// Now has junk data
printf("Destroyed: %d\n", fp->f_count);

exit(EXIT_SUCCESS);
}
```